

# Chapter 4. How to Create Reusable Infrastructure with Terraform Modules

---

At the end of [Chapter 3](#), you deployed the architecture shown in [Figure 4-1](#).

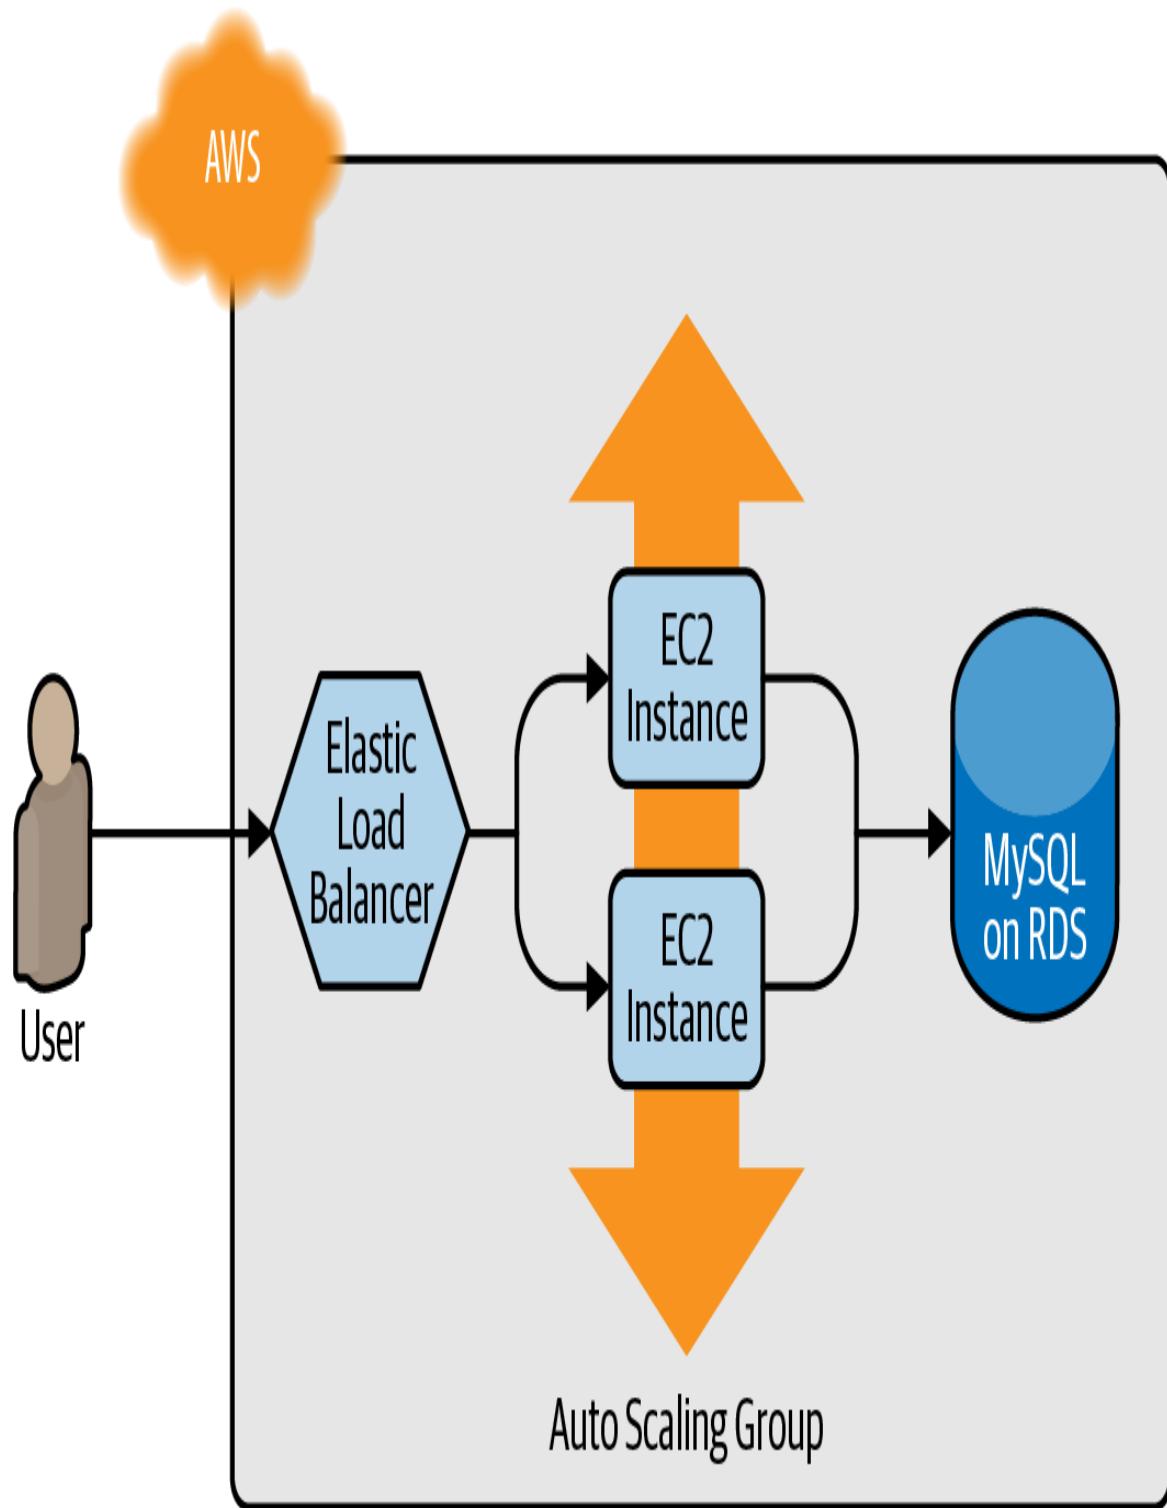
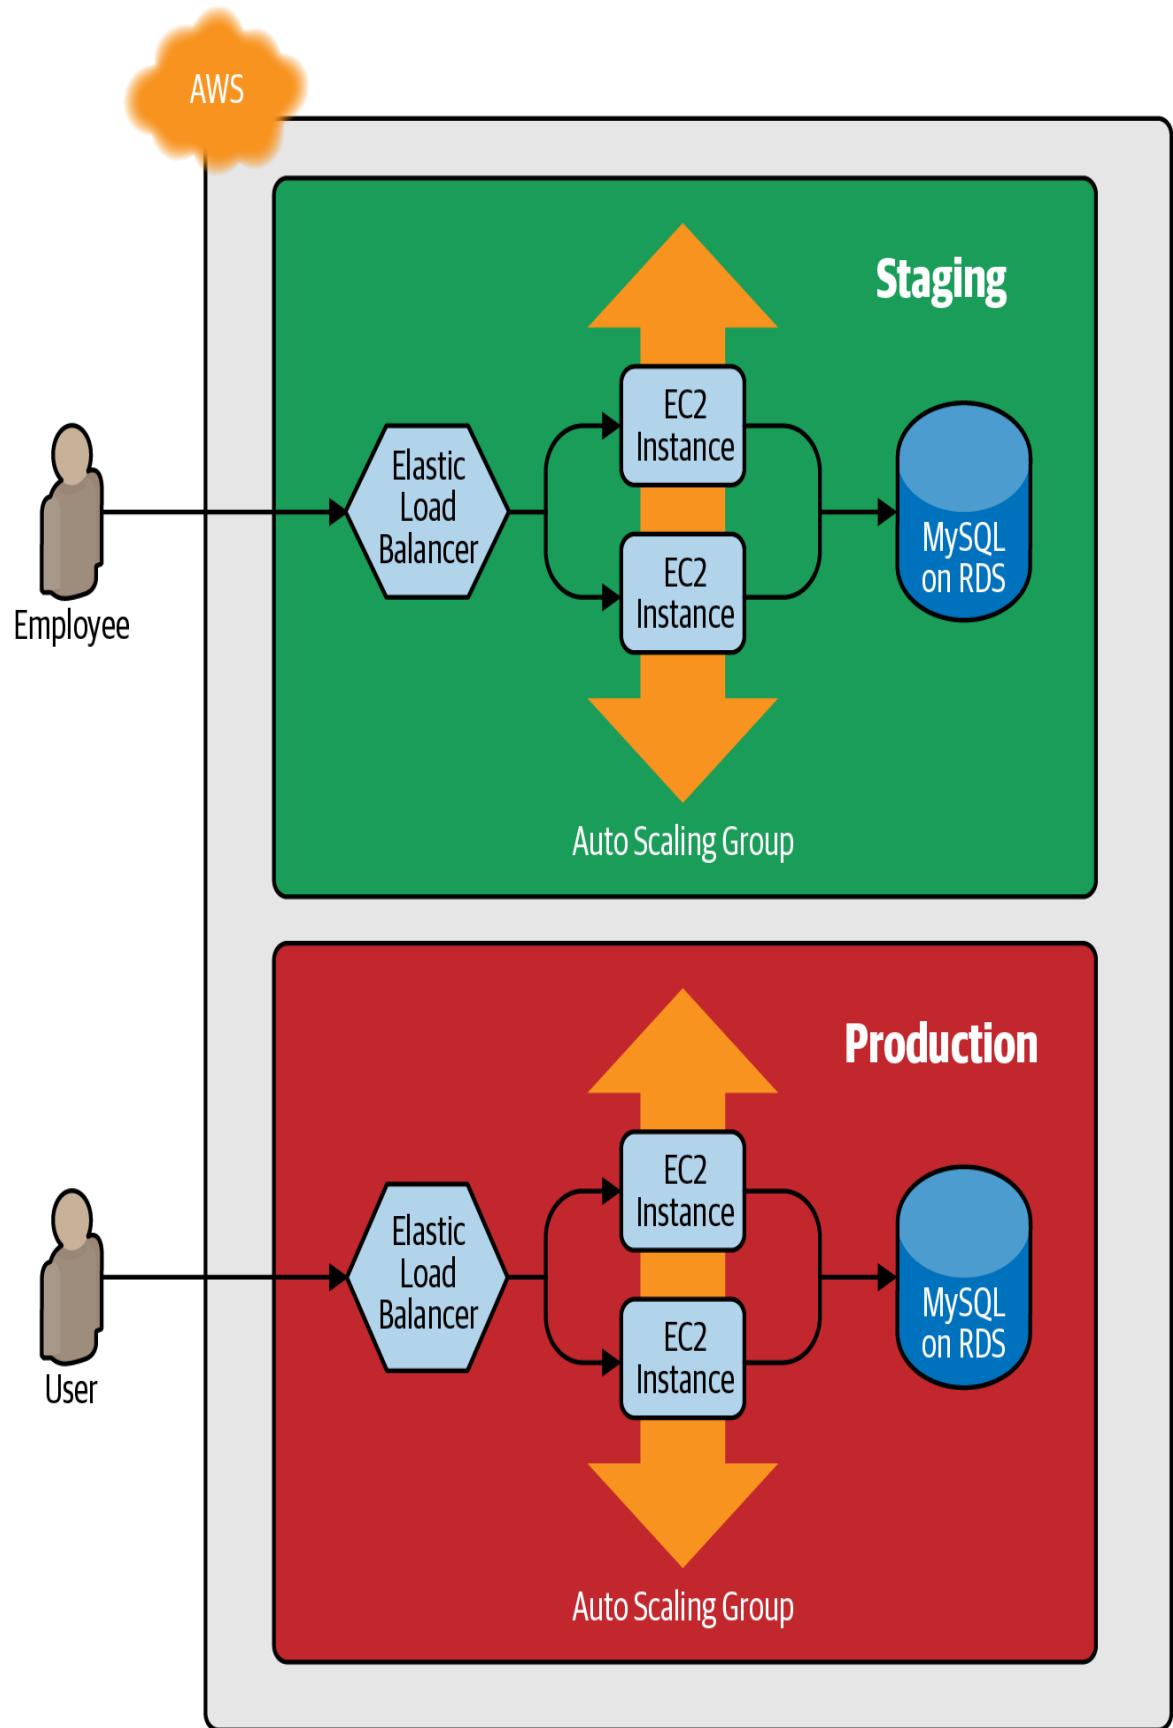


Figure 4-1. The architecture you deployed in previous chapters included a load balancer, web server cluster, and database.

This works great as a first environment, but you typically need at least two environments: one for your team's internal testing ("staging") and one that

real users can access (“production”), as shown in [Figure 4-2](#). Ideally, the two environments are nearly identical, though you might run slightly fewer/smaller servers in staging to save money.



*Figure 4-2. The architecture you'll deploy in this chapter will have two environments, each with its own load balancer, web server cluster, and database.*

How do you add this production environment without having to copy and paste all of the code from staging? For example, how do you avoid having to copy and paste all the code in *stage/services/webserver-cluster* into *prod/services/webserver-cluster* and all the code in *stage/data-stores/mysql* into *prod/data-stores/mysql*?

In a general-purpose programming language such as Ruby, if you had the same code copied and pasted in several places, you could put that code inside of a function and reuse that function everywhere:

```
# Define the function in one place
def example_function()
    puts "Hello, World"
end

# Use the function in multiple other places
example_function()
```

With Terraform, you can put your code inside of a *Terraform module* and reuse that module in multiple places throughout your code. Instead of having the same code copied and pasted in the staging and production environments, you'll be able to have both environments reuse code from the same module, as shown in [Figure 4-3](#).

